

Scalability Limits of HDFS

Christopher Chute, David Brandfonbrener, Leo Shimonaka, Matthew Vasseur

May 2, 2016

Abstract

1 Introduction

2 HDFS Architecture

Hadoop claims to be built on five core principles that govern the architectural decisions. These are (1) hardware failure is the norm. As a result, HDFS wants to be highly reliable in the face of hardware failure so the architecture incorporates replication of data across multiple nodes. Then HDFS wants to (2) allow streaming data access and (3) support very large data sets. With these goals in mind, HDFS wants to use a distributed system to accommodate large files and implement distributed reads to prevent streaming access from clogging the system. Next HDFS wants to (4) provide simple concurrency control with a write-once read-many model. This is less a feature and more a decision to weigh reads over writes by providing highly available and concurrent reads at the expense of allowing concurrent writes. And lastly, HDFS wants (5) portability, which makes sense as HDFS must work on different computer architectures

To satisfy these principles, the HDFS architecture was designed as follows.

2.1 Namenode

2.2 Datanode

3 Limitations of Single Namenode Architecture

3.1 Physical Memory Size

3.2 Namenode CPU Bottleneck

4 Performance Evaluation

4.1 Experiment 1: Memory Experiments

4.2 Experiment 2: CPU Experiments

4.3 Analysis

5 Improvements to Single Namenode Architecture

6 Conclusion